SUBMISSION 1:

How many rows are missing a value in the "State" column? Explain how you came up with the number.

It appears that **5377** rows return “Blank” for their state.

I came up by my answer by looking at the text facet of the state column. See left panel in below picture. I also sorted the data by the blank fields to double check they were blank.

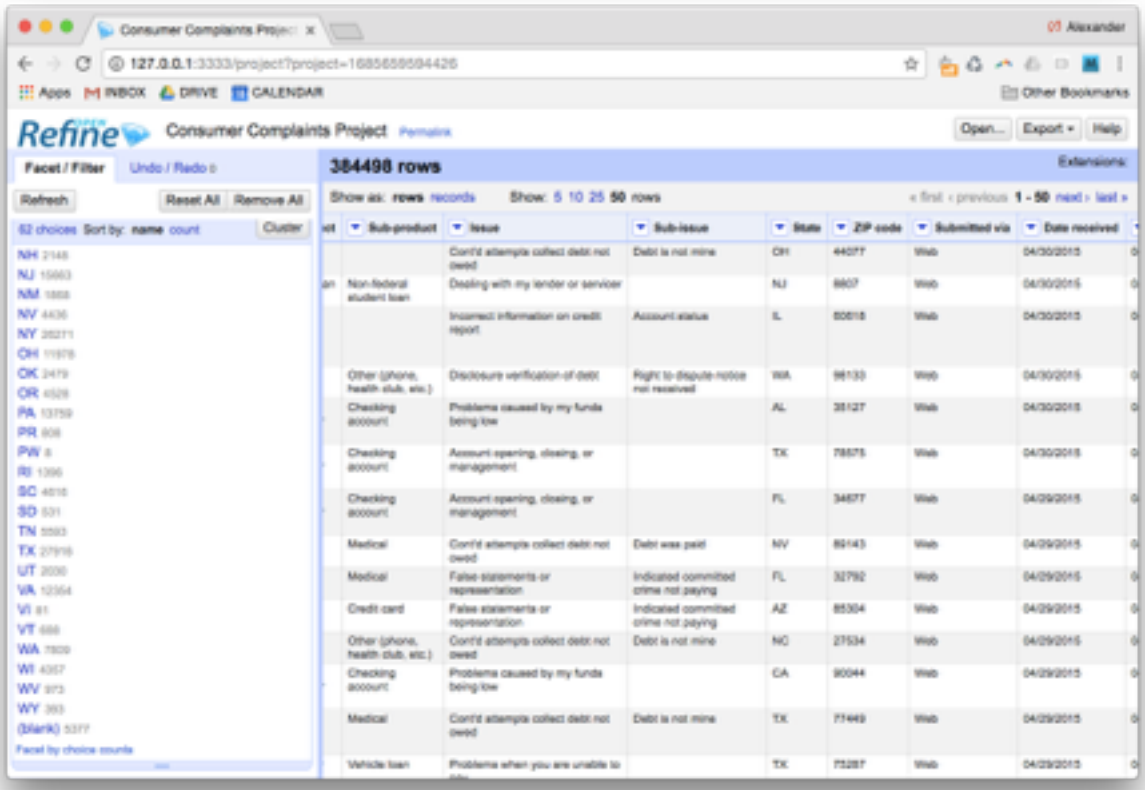

SUBMISSION 2:

How many rows with missing ZIP codes do you have?

It appears that **4362** rows are blank. This is *after* transforming the zip code text into number format.

SUBMISSION 3:

If you consider all ZIP codes less than 99999 to be valid, how many valid and invalid ZIP codes do you have, respectively?

I don't completely understand this question but here's what I think it's saying:

```
if zip code == 99999:
        zip code is INVALID

if zip code < 99999:
        zip code is VALID

if zip code == "Blank"
        zip code is INVALID

Total INVALID = count(99999) + count(blank)
                        =  34961 + 4362
                        =  39323
```

So, total invalid is **39,323**. Total valid would be **345,175** (total rows *minus* total invalid or 384498 - 39323).

SUBMISSION 4:

Change the radius to 3.0. What happens? Do you want to merge any of the resulting matches?

When we changed the radius to 3.0, the cluster began to pair words that were similar but were actually distinct.

So, for example, Indonesia and Micronesia are NOT misspellings but rather two different regions with very similar spelling.

You don't want to merge these particular values.

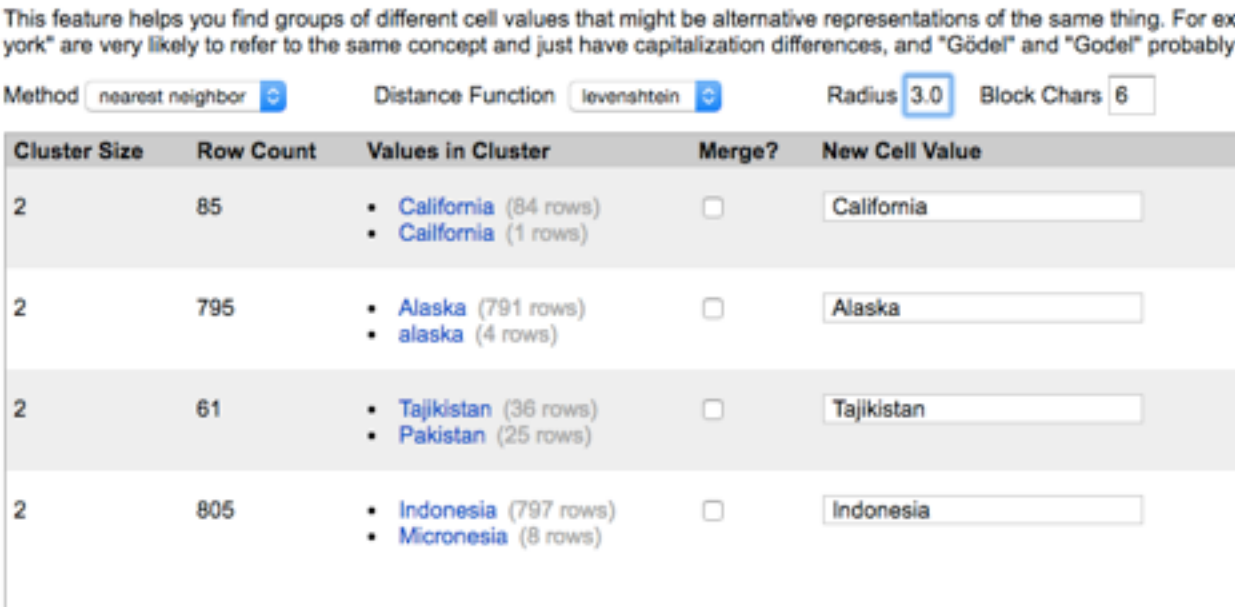

SUBMISSION 5:

Change the block size to 2. Give two examples of new clusters that may be worth merging.

First: Two clear misspellings of California:

Cluster 1

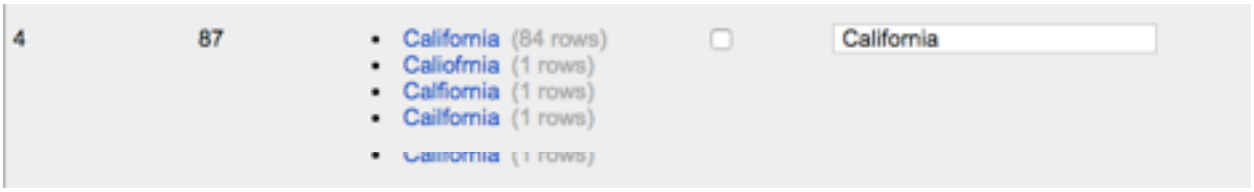

Cluster 2

Second: Two groupings that seem to be primarily Alaska. This would distort the data by folding in entries from 3 other regions. It was the “best worst” option presented in the group.

SUBMISSION 6:

Explain in words what happens when you cluster the "place" column, and why you think that happened. What additional functionality could OpenRefine provide to possibly deal with the situation?

The cluster seems to take a long time to process. Waited a few minutes before canceling.

This could have happened because:

a. There were too many individuals terms to process
b. The terms were different data types
c. There was too much data to process
d. The clusters were so small and so populous that it was running for a long time before completing

OpenRefine could provide a few things:

a. Split function that groups similar data types
b. A lighter weight visualization tool that just shows you all the unclustered data color coded by type
c. General rules about ideal data sets to cluster

SUBMISSION 7:

Submit a representation of the resulting matrix from the Levenshtein edit distance calculation. The resulting value should be correct.

```
>>> from Levenshtein import *
>>> distance("hej","hei")
1
>>> distance("monthgomery st","montgomery street")
5
>>> distance("gumbarrel","gunbarell")
3
>>>
```

		1	2	3	4	5	6	7	8	9	10
			g	u	m	b	a	r	r	e	l
1		0	1	2	3	4	5	6	7	8	9
2	g	1	0	1	2	3	4	5	6	7	8
3	u	2	1	0	1	2	3	4	5	6	7
4	n	3	2	1	1	2	3	4	5	6	7
5	b	4	3	2	2	1	2	3	4	5	6
6	a	5	4	3	3	2	1	2	3	4	5
7	r	6	5	4	4	3	2	1	2	3	4
8	e	7	6	5	5	4	3	2	2	2	3
9	l	8	7	6	6	5	4	3	3	3	2
10	l	9	8	7	7	6	5	4	4	4	3